

AI Agents in an Experimental Technical Services Workflow

Jeremy Nelson, Stanford University Libraries

DISCLAIMER: Generative AI was utilized in the creation of this presentation

What are AI Agents?



They act, not just answer — chatbots return text; agents do the work.



Tools are their hands — APIs, scripts, and commands define what they can reach.



They work in a loop — observe, decide, act, check the result, repeat.



Instructions, not code — you describe the goal in plain language; the agent works out the steps.



Humans stay accountable — people approve the decisions that matter, and can always stop the run.

AI Agentic Harnesses



The harness is the runtime

Around the model. It holds the loop, tools, and permissions.



What they all provide

File access, shell, API & MCP tools, skills, approval gates.



Commercial

Vendor-tuned, tied to their own models: [Claude Code](#), [OpenAI Codex](#).



Open source

Model-agnostic, self-hostable: [OpenCode](#), [goose](#).



Why it matters

Commercial buys capability; open source buys control.



Portable

The workflow is markdown and MCP, not vendor lock-in.

Using AGENTS.md



A README for the agent

how to set up, run, and work in the project.



One file, many harnesses

no required structure; just markdown the agent parses.



Best practice: be specific and correct

exact commands, real file paths, explicit prohibitions ("never commit `.env`"). A stale instruction is worse than none.



Best practice: treat it as code

commit it, review it, update it when the project moves.



In this project

it names the environment variables, the three-stage skill pipeline, and where the authoritative BIBFRAME mappings live.

About SKILL.md



A skill is a folder with a SKILL.md

`name` and `description` frontmatter, then instructions.



Progressive disclosures

agents load only the name and description at startup, then the full instructions when a task matches.



Open format

write a skill once, run it in any skills-compatible agent.



Best practice: the description is the trigger

say *when* to use the skill, not just what it does.



Best practice: keep SKILL.md lean

detail goes in `references/`, deterministic steps go in `scripts/`

Example Technical Services Workflow



folio-inventory-retrieval SKILL.md



Frontmatter Trigger

A **name** and a one-line **description** are all the agent needs to know when to reach for this skill.



Instructions, Not an App

Hands the agent a **FolioClient** snippet, credentials from **.env**, and links to Inventory and Search APIs. No wrapper code.



Three Ways In, One Outcome

UUID or HRID goes straight to Inventory. Otherwise, search by title/author as a fallback and ask user to choose before pulling data.

Demos with Claude Code & Open Code

01. Inspection

Show the 36-line **SKILL.md** — the entire stage, detailed in clear prose.

02. Direct Retrieval

Run with just an HRID; watch the agent authenticate and fetch Instance + Holdings JSON.

03. Search Fallback

Run with only title and author, trigger the interactive cataloger selection prompt.

bibframe-transform SKILL.md



The Mapping Table is the Code

Two markdown tables carry FOLIO Instance and Holdings JSON into BIBFRAME Work and Instance JSON-LD, field by field. No transformation script to maintain.



Cataloging Judgment, Written Down

FOLIO reference-data UUIDs resolve to labels through the same `FolioClient`; unlinked controlled values become blank nodes with `rdfs:label`, not invented URIs.



Hard-Won Constraints

Don't set `bf:hasInstance` before the Work exists, and don't dedupe identifiers that share a value under different types. Quirks are documented, not rediscovered.

Demo with Claude Code & Open Code

01. Inspection

Show the `SKILL.md` — prerequisites, JSON-LD conventions, and the two mapping tables.

02. Direct Transformation

Hand the agent the Instance JSON; watch it resolve reference data and emit Work + Instance JSON-LD.

03. Auditable Mapping

Point at one table row and its triple — mapping is auditable by a cataloger, not just a developer.

04. Expected Scaling

~30 barcodes × 3 identifier types = ~90 `bf:identifiedBy` entries. Expected behavior, not a bug.

bibframe-validation SKILL.md

blue-core-mcp-ingestion SKILL.md

Demonstration

Thank-you!

Email: jpnelson@stanford.edu

Source code: <https://github.com/jermnelson/tech-services-ai-workflow>