




Beyond the Chatbot: Agentic AI and the Future of Library Workflows

Jeremy Nelson, Taryn Marks, Will
Huggins

Sept. 1, 11am



Learning Objectives




- Today, you will:
 - develop a mental model for identifying agentic AI opportunities in your own institutions;
 - gain familiarity with current platforms and approaches for building agentic workflows, along with a framework for evaluating the security, governance, and appropriateness of deploying AI agents; and
 - consider the implications for widespread library collaborations

Agenda



- Overview: framework for identifying agentic opportunities, and evaluating the security, governance, and appropriateness of deploying AI agents
- Demos: Applying the framework to specific demos
 - using different platforms & approaches
- Assessing broader implications

The background is a dark, blue-tinted aerial photograph of a city, likely Prague, featuring a large bridge with multiple arches spanning a river. The city's architecture and hills are visible in the background, creating a scenic backdrop for the title.

Agentic Framework for Librarians

SCOPE: An Overview



❏ SCOPE

- ❏ Suitable?
- ❏ Content?
- ❏ Oversight?
- ❏ Permissions?
- ❏ Evaluation?

SCOPE: An Agentic AI Framework for Libraries



- ❑ S: Suitable:
 - ❑ Is the task suitable for an agent? Can you design an agent that's suitable for the specific task?
- ❑ C: Content:
 - ❑ What information does the agent need (both to create it and to deploy it), and can you give it that safely?

SCOPE: An Agentic AI Framework for Libraries



- ❑ O: Oversight
 - ❑ Who will be responsible for the initial and ongoing evaluation of governance and security? Where should you build in explicit human checkpoints?
- ❑ P: Permission
 - ❑ What permissions do you need (internal, external, others)?
- ❑ E: Evaluate
 - ❑ How will you assess the agent as it performs the task? How will you assess any output? What will you do if it fails?

Examining Tasks for Agentic Possibilities



❑ Literature review, research workflow

❑ Suitable?

- ✓ Workflow-based; repeatable; verifiable
- ✗ Clear goals; professional judgment; conclusions from research

❑ Content?

- ✓ Research workflow; connectors; best databases; result format; guardrails
- ✗ Confidential, private data

❑ Oversight?

- ✓ Policies; human checkpoints; external actions; built into design
- ✗ Who owns; regular assessment

Examining Tasks for Agentic Possibilities



❏ Research workflow: The Literature Review

❏ Permissions?

- ✓ Connectors, APIs, institution, IT, databases, copyright
- ✗ Paywalls, cost

❏ Evaluation?

- ✓ Built-in verification/assessment of sources and claims; search rigor; human checkpoints; output evaluation
- ✗ Agent's actual performance? Responsibility for mistakes?

SCOPE: Conclusion



❏ SCOPE

- ❏ Suitable?
- ❏ Content?
- ❏ Oversight?
- ❏ Permissions?
- ❏ Evaluation?



Agentic Tools

An Agent is a Stack—Not a Single Product



A product name tells us where the work is coordinated.

It does not tell us what the agent can access—or what it is allowed to do.

THE AGENTIC WORKFLOW STACK

CONTROLS

Permissions · approvals · logs

AGENT HARNESS

Plans and coordinates

TOOLS

APIs · MCP · files · code

CONTEXT

Instructions · records · policies

MODEL

Interprets and generates

Three Build Paths Trade Speed for Control



Automation follows a path. An agent can choose among the paths you approve.

EXPLORE

Sandbox

Learn in a front-facing AI workspace

ChatGPT · Claude

EXTEND

Assemble

Combine reusable instructions with connected tools

Skills · plugins · connectors

BUILD

Customize

Work directly with files, APIs, and MCP

Claude Code · Codex Agent SDKs

Faster start

More control—and more technical ownership


Use ChatGPT as a Sandbox for Learning Agents



Start with one bounded, non-sensitive task



Skill = how to work | Connector = what it can reach | Agent = coordinates the work



Demo 1: Research Workflow Agent (Pre-Recorded)

skills-excerpt.vercel.app



APIs and MCPs Play Different Roles



API

A system-specific doorway

Defines the requests another application can make—and the responses the system returns.

MCP

An AI-facing tool layer

Permits an AI host discover and invoke defined tools or retrieve contextual resources.

Often used together: MCP can expose an API operation as an AI-usable tool.

Today's Demos Combines Agentic and Deterministic Work



Goal: turn a FOLIO inventory record into linked-data resources



Agentic planning | Structured output | Controlled write




Demo 2: Technical Services Pipeline

Details on Technical Services Pipeline Demo




- Demonstration of an example workflow that spans different systems
- Blue Core Project - shared BIBFRAME data licensed under Open Metadata License, currently developing MVP but has existing prototype. Details at <https://bluecore.info/>
- Source code for demo available at <https://github.com/jermnelson/tech-services-ai-workflow/>



Demo 3: Institutional Knowledge Assistant

Stanford University - Admin Search Guide



- Part of Stanford's AI Playground
- RAG (Retrieval Augmented Generation) of existing guidelines that govern workplace interactions, approaches, procedures, and processes at <https://adminguide.stanford.edu/>
- Available to authenticated Stanford employees



WOLFcon 2026

Prague, Czech Republic

Aug. 31 - Sept. 3

Thank You!

Contact us!

- Jeremy Nelson,
jpnelson@stanford.edu
- Will Huggins
- Taryn Marks, tlmarks@stanford.edu